

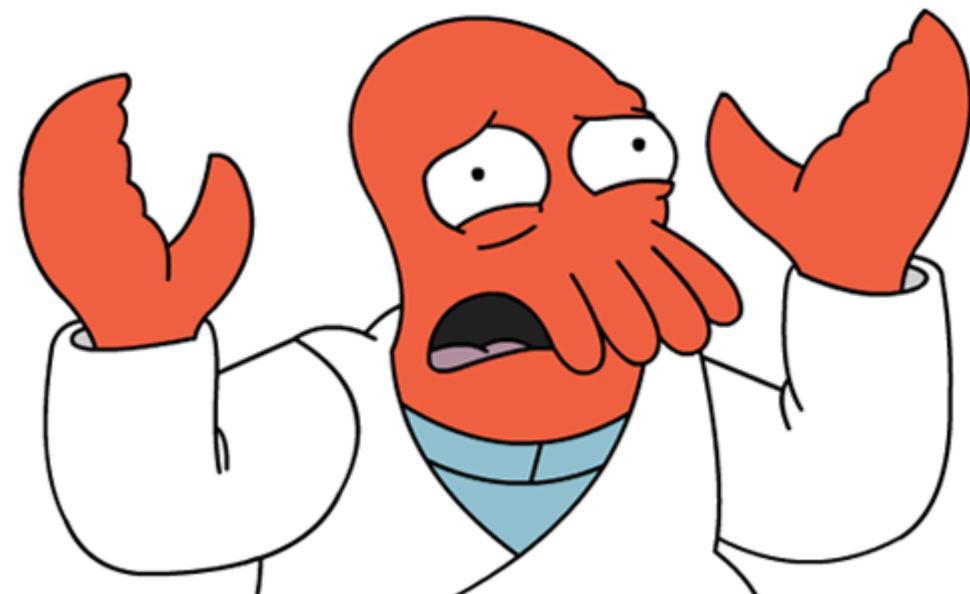
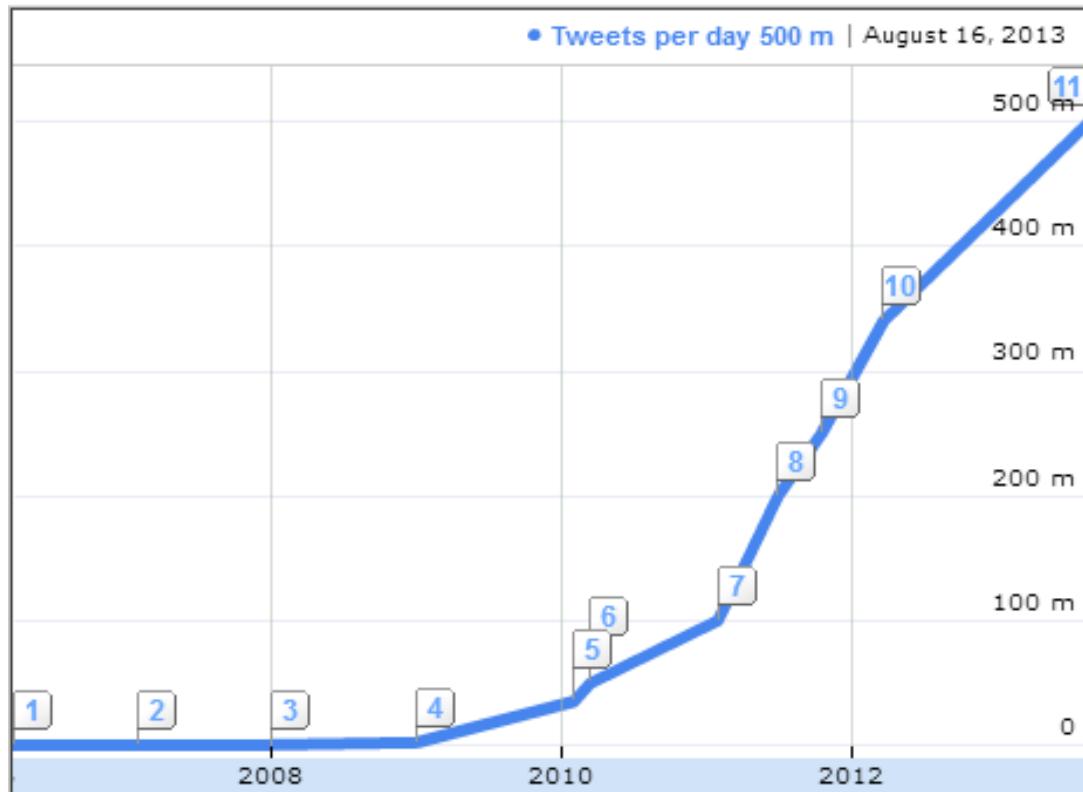
# Stream Processing

---

DAVID OSTROVSKY | PROOFPOINT

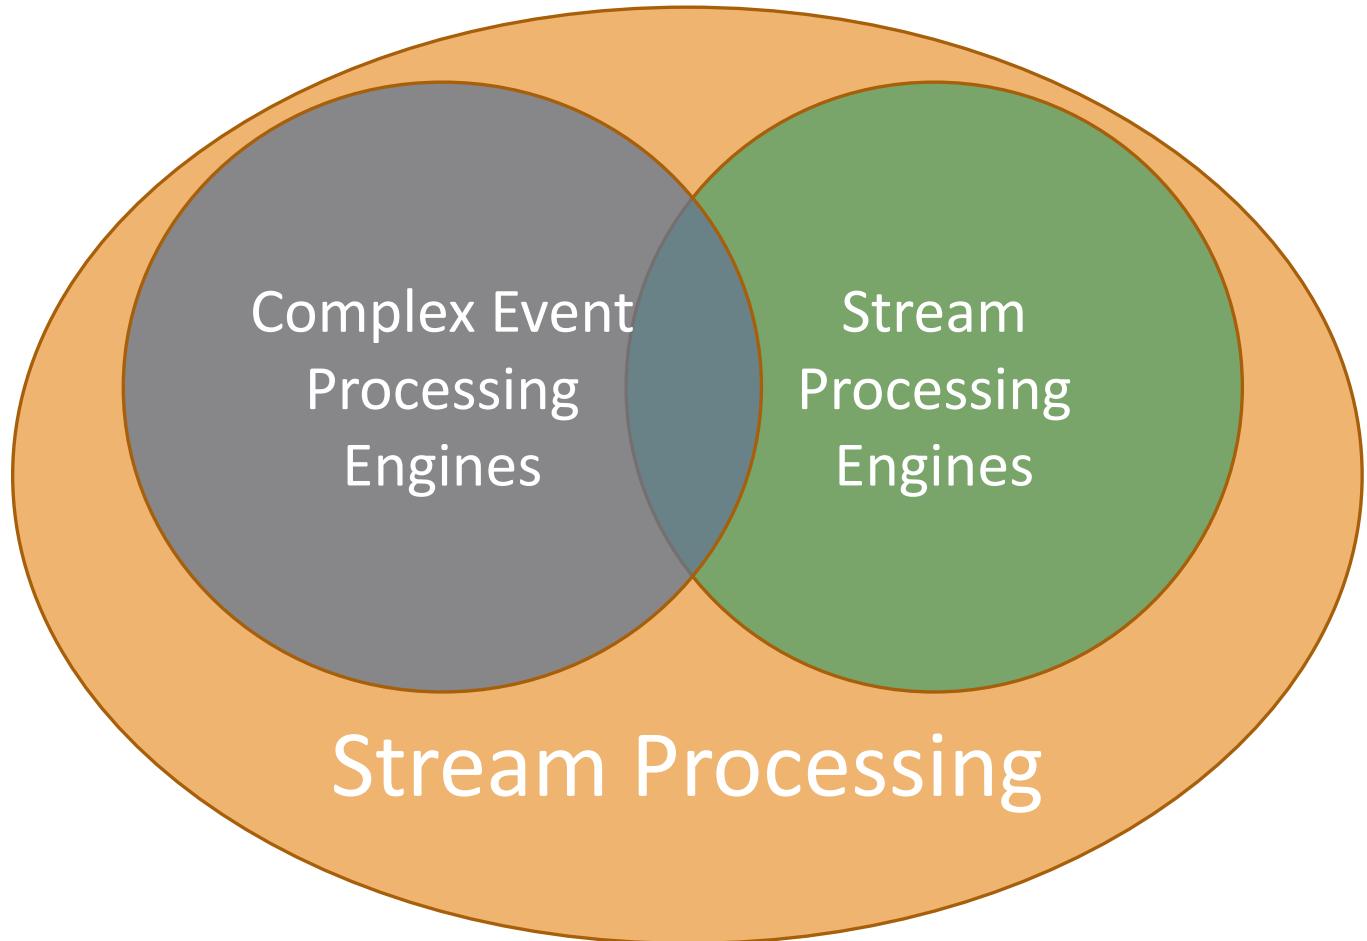
# Why Streaming?

---

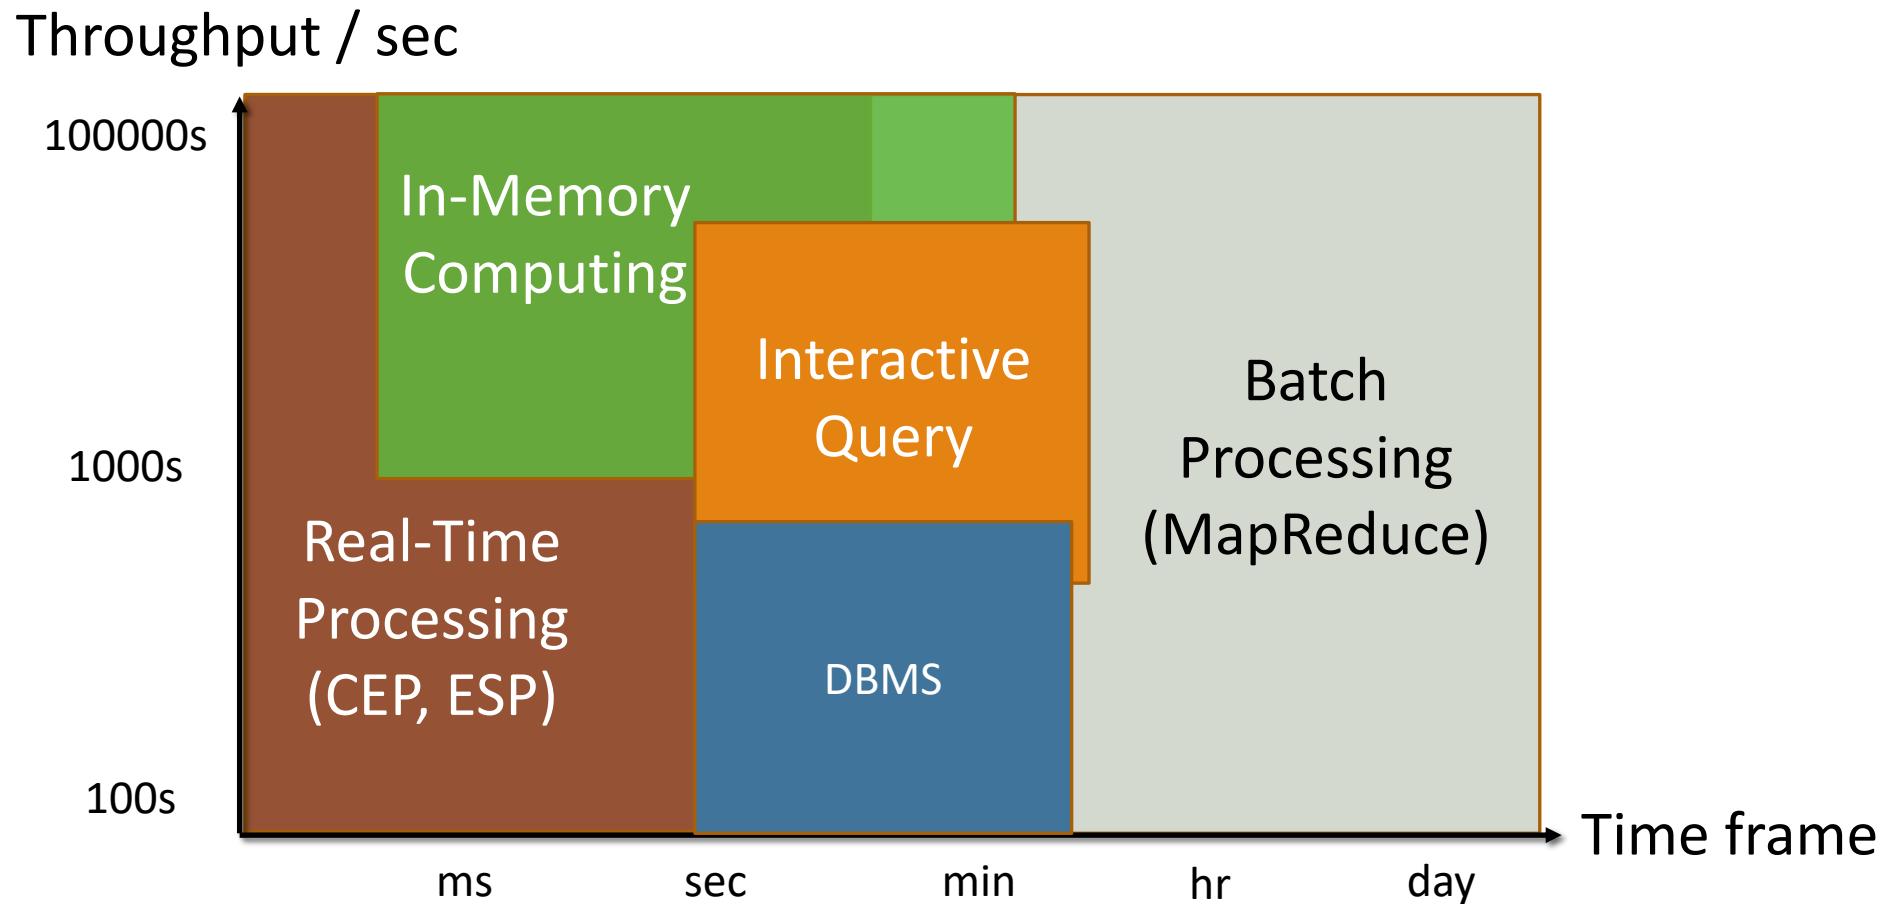


# Streaming Data

---



# Types of Data Processing



# All Apache, all the Time

---



# No Love for Microsoft?

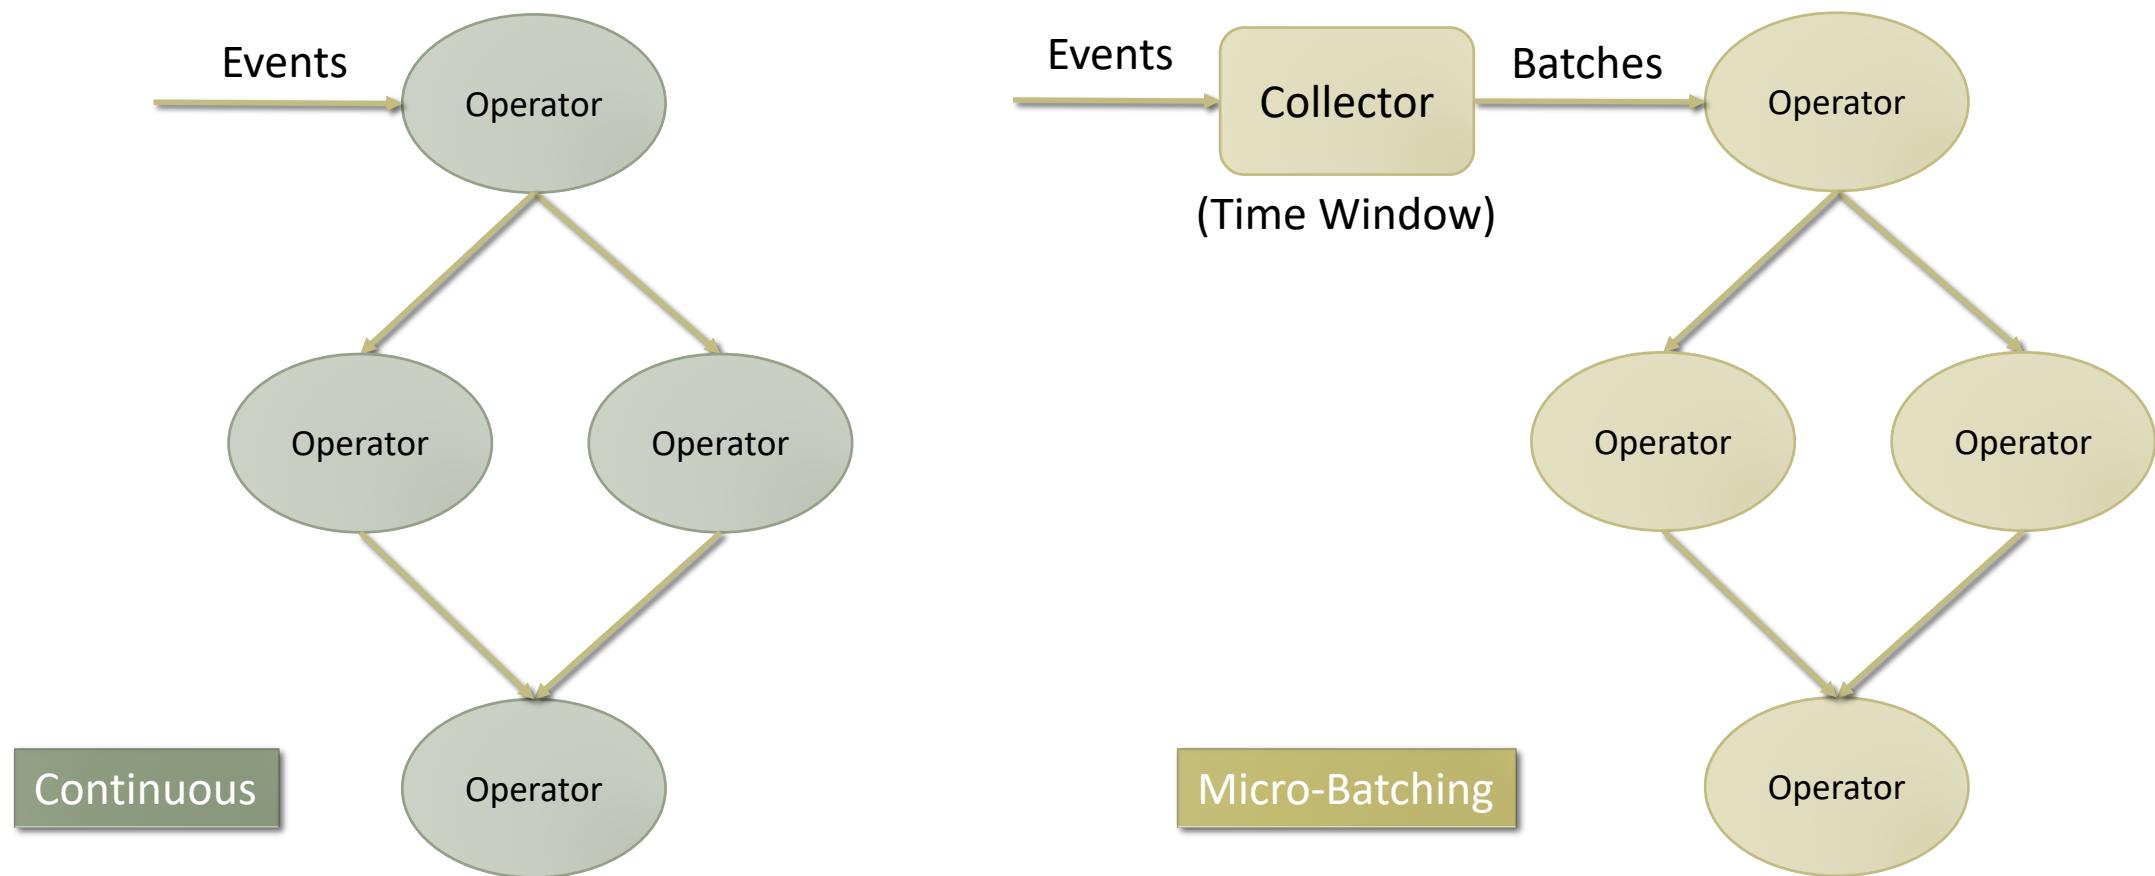
---

# Orleans



# Processing Model

---



# Programming Model

---



**Continuous**



**Micro-Batch**



**Micro-Batch**



**Continuous**



**Continuous\***

\* Has a batch abstraction on top of streaming

# API and Expressiveness

---

## Compositional

```
public class PrinterBolt extends BaseBasicBolt  
{  
  
    public void execute(Tuple tuple, ...) {  
        System.out.println(tuple);  
    }  
}
```

```
topology.setBolt("print", new PrinterBolt())  
    .shuffleGrouping("twitter");
```

## Declarative

```
val ssc = new StreamingContext(conf, Seconds(1))  
ssc.socketTextStream("localhost", 9999)  
    .flatMap(_.split(" "))  
    .map(word => (word, 1))  
    .reduceByKey(_ + _)  
    .print()
```

# API and Expressiveness

---



Compositional	Compositional	Declarative	Both	Declarative
JVM, Python, Ruby, JS, Perl	JVM	JVM, Python	JVM	JVM, Python*

\* Only for the DataSet API (batch)

# Storm + Trident

---

Topology:

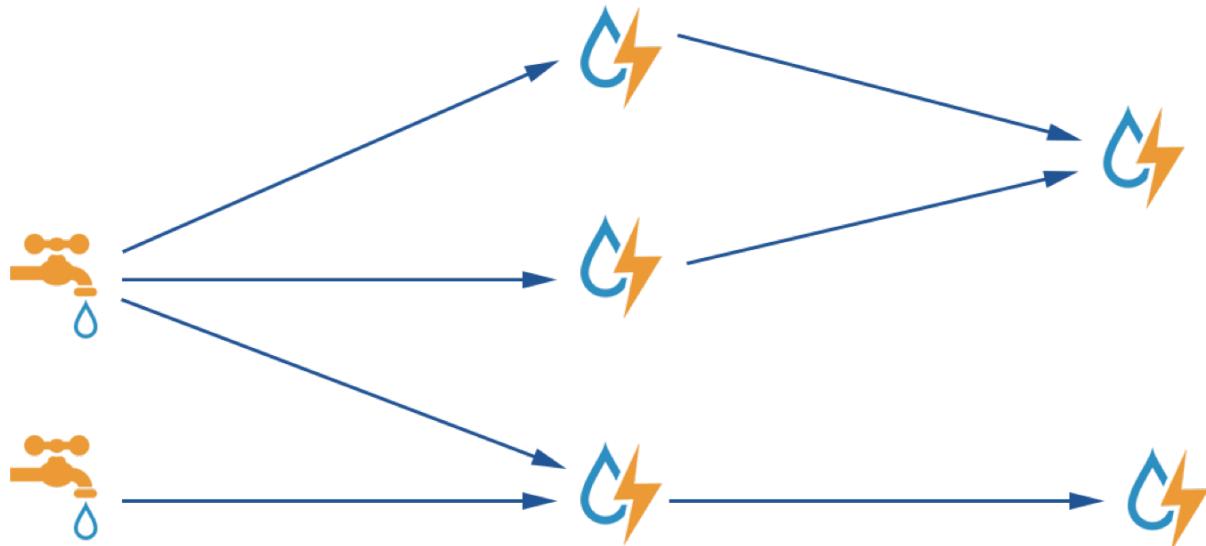
- Spouts
- Bolts

Stream Groupings:

- Shuffle
- Fields
- All
- ...

Nimbus (Master)

- Workers



# Spark Streaming

---

Resilient Distributed Datasets (RDD)

DStreams – sequences of RDDs



# Kafka Streams

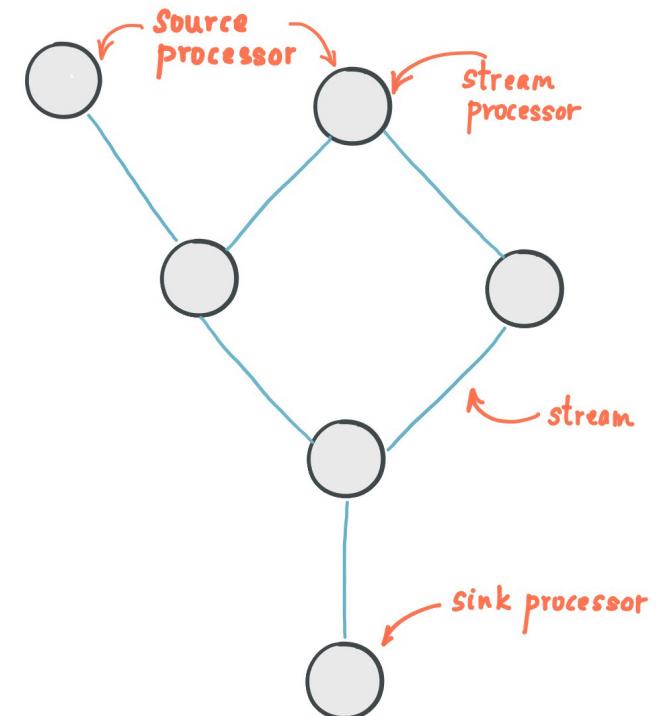
---

Uses Kafka for streaming

- Topics (streams)
  - Partitioned across Brokers
- Producers
- Consumers

Applications use the Streams API

- KStream
- KTable



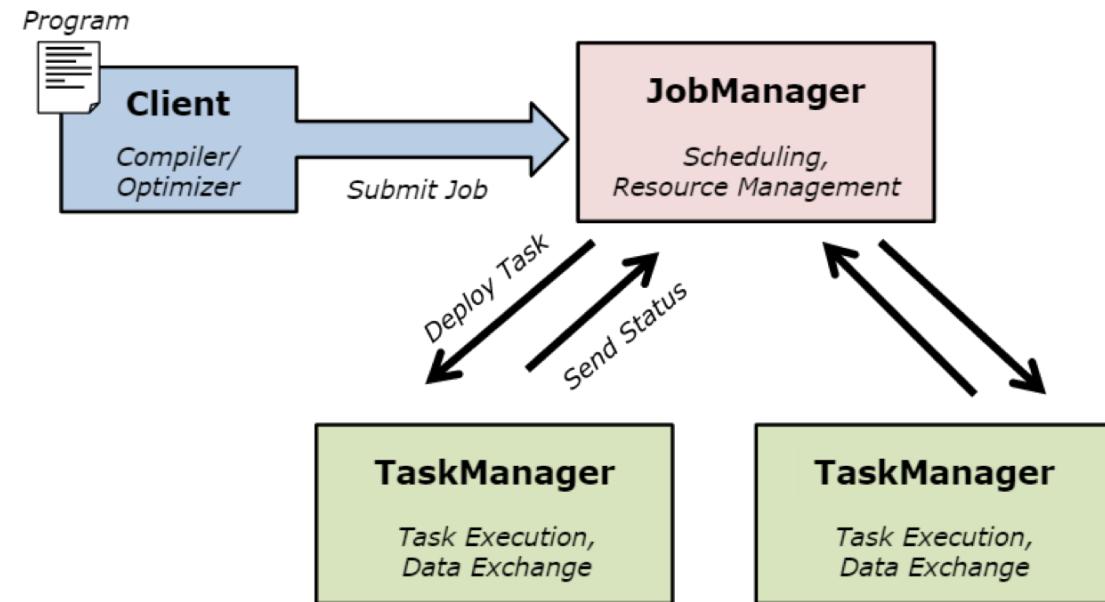
PROCESSOR TOPOLOGY

# Flink

---

## Dataflows

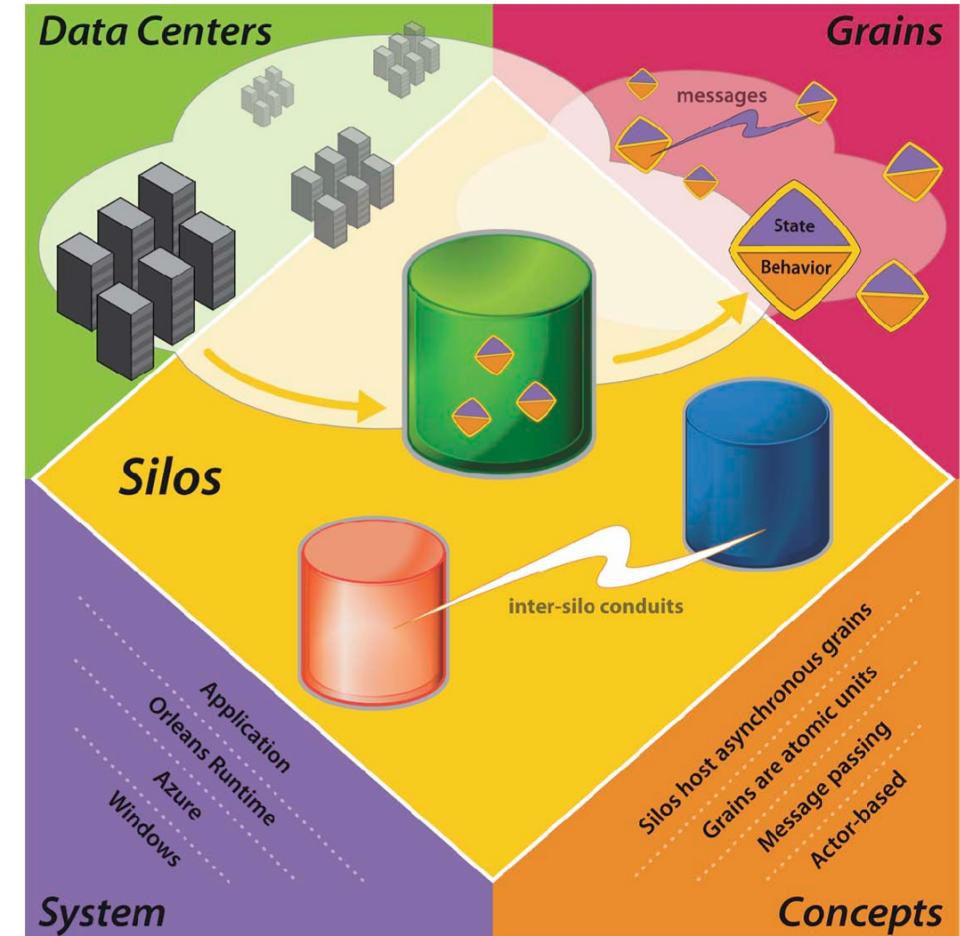
- Streams
  - Source(s)
  - Sink(s)
- Transformations (operators)



# Orleans

## Virtual Actor System in .NET

- Grains (operators)
- Silos (containers)
- Streams



# Message Delivery Guarantees

---

At Most Once

At Least Once

Exactly Once

<b>Source</b>	Sockets Twitter Streaming API Any non-repeatable	Files Simple Queues Any forward-only	Kafka, RabbitMQ Collections Stateful
<b>Sink</b>		Data Stores Sockets Files	HDFS rolling sink

# Highest Possible Guarantee

---



Spark  
Streaming



---

**At least once\***

**Exactly once\***

**Exactly once\*\***

**Exactly once\***

**Exactly once\***

---

\* Doesn't apply to side-effects

\*\* Only at the batch level

# Reliability and Fault Tolerance

---



---

ACK per tuple

RDD checkpoints

Partition offset  
checkpoints

Barrier  
checkpoints

---

# State Management

---



---

**Manual**

**Dedicated state providers (memory, external)**

**RDD with per-key state**

**Local K/V store + changelog in Kafka**

**Stored with snapshots, configurable backends**

---

# Performance

---



Spark  
Streaming



<b>Latency</b>	Low	Medium	Medium-High*	Low	Low**
<b>Throughput</b>	Medium	Medium	High	High	High

\* Depends on batching

\*\* For streaming, not micro-batching

# Extended Ecosystem

---



**SAMOA (ML)**



**Trident-ML**



**Spark SQL,  
MLlib  
GraphX**



**N/A**



**CEP  
Gelly\*  
FlinkML\*  
Table API (SQL)\***

---

\* DataSet API (batch)

# Production and Maturity

---



---

**Mature,  
many users,  
260 contributors**

**Relatively mature,  
many users  
1092 contributors\***

**New,  
built on mature  
components,  
enterprise backing,  
306 contributors**

**New,  
high momentum,  
fewer users,  
321 contributors**

---

\* Spark, not just spark streaming

\*\* Contributor numbers as of 12/6/2017



*Everybody got that?*